Questions an audience would ask

1. Is the accuracy score of the model accurate enough?
    a. Based on the results, the model can be used to make an initial assessment of a brain stroke in a patient. Over time, as variables are updated and more variables are added, the accuracy will improve.
2. Is there a way to prove the model is working as intended?
    a. On a periodic basis, the diagnosis of a brain stroke will be sampled to confirm the accuracy.
3. How can the model be improved?
    a. The model can be improved by updating variables and records and making sure they are medically correct. In addition, other variables than can affect a brain stroke should also be added.
4. What other models can be built that might have better results?
    a. Some other models that can be built to compare the results with are Naive Bayes and K-Nearest Neighbours.
5. What variables have the biggest effect on brain stroke?
    a. The initial heatmap showed that the age of the patient had the biggest effect on a stroke.
6. Does the data violate a person's privacy?
    a. Looking at the dataset there are no specific identifiers to find the identity of the patient.
7. What variables have the least effect on brain stroke?

   a. Based on the initial heatmap, it looks like the bmi had the least effect on brain stroke.

8. Can the model be used as it is?
   a. I believe the model and dataset should be looked at by healthcare professionals to confirm the accuracy of the records. More variables should be added to increase the accuracy as well.
9. What if it inaccurately predicts a person having a brain stroke?
   a. The model and dataset will be looked at to improve the accuracy.
10. What is the point of this model when doctors can give a diagnosis?
   a. The point of this model is to help doctors make a quicker and more accurate decision when diagnosing a patient with a brain stroke.